

Homework 3

Gonzalo Heras 103401
Tomó 4node 103239

① a)

	y_1	y_2	t
x_1	0,7	-0,3	0,8
x_2	0,4	0,5	0,6
x_3	-0,2	0,8	0,3
x_4	-0,4	0,3	0,3

$$C = \left[\begin{pmatrix} 0 \\ 0 \end{pmatrix}; \begin{pmatrix} 1 \\ -1 \end{pmatrix}; \begin{pmatrix} -1 \\ 1 \end{pmatrix} \right]$$

$C_1 \quad C_2 \quad C_3$

$0,1$

$$\eta = 1$$

$$\phi_j(x) = \exp \left(- \frac{\|x - C_j\|^2}{2} \right)$$

$$X = \begin{bmatrix} 1 & \phi_1(x_1) & \phi_2(x_1) & \phi_3(x_1) \\ 1 & \phi_1(x_2) & \phi_2(x_2) & \phi_3(x_2) \\ 1 & \phi_1(x_3) & \phi_2(x_3) & \phi_3(x_3) \\ 1 & \phi_1(x_4) & \phi_2(x_4) & \phi_3(x_4) \end{bmatrix} = \begin{bmatrix} 1 & 0,74826 & 0,74826 & 0,10107 \\ 1 & 0,81465 & 0,27117 & 0,33121 \\ 1 & 0,71177 & 0,09633 & 0,71177 \\ 1 & 0,88250 & 0,16122 & 0,65377 \end{bmatrix}$$

$$W = (X^T \cdot X + \lambda \cdot I)^{-1} \cdot X^T \cdot t = \begin{bmatrix} 0,33914 \\ 0,19945 \\ 0,40096 \\ -0,29600 \end{bmatrix}$$

equação do hiperplano: $w_0 + w_1 \cdot y_1 + w_2 \cdot y_2 + w_3 \cdot y_3 = \hat{z}$

→ $0,33914 + 0,19945 \cdot y_1 + 0,40096 \cdot y_2 - 0,29600 \cdot y_3 = \hat{z}$

①

b)

$$\hat{t} = X \cdot W = \begin{bmatrix} 0,75844 \\ 0,51232 \\ 0,30905 \\ 0,38629 \end{bmatrix}$$

$$RMSE = \sqrt{\frac{1}{m} \sum_{i=1}^m (t_i - \hat{t}_i)^2} = 0,06508$$

(2)

2)

$$w^{[1]} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 2 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix} \quad w^{[2]} = \begin{bmatrix} 1 & 4 & 1 \\ 1 & 1 & 1 \end{bmatrix} \quad w^{[3]} = \begin{bmatrix} 1 & 1 \\ 3 & 1 \\ 1 & 1 \end{bmatrix}$$

$$b^{[1]} = [1 \ 1 \ 1]^T$$

$$b^{[2]} = [1 \ 1]^T$$

$$b^{[3]} = [1 \ 1 \ 1]^T$$

$$f(x) = \tanh(0,5 \cdot x - 2) \quad \eta = 0,1$$

$$E(w) = \frac{1}{2} \|z - \hat{z}\|^2$$

Nota: Todos os cálculos foram realizados com auxílio da biblioteca numpy.

$$\rightarrow x_1 = [1 \ 1 \ 1 \ 1]^T, t_1 = B = [-1 \ 1 \ -1]^T$$

$$\rightarrow x_2 = [1 \ 0 \ 0 \ -1]^T, t_2 = A = [1 \ 0 \ 0 \ -1]^T$$

Forward Propagation $\rightarrow x^{[p]} = \phi^{[p]} \cdot (w^{[p]} x^{[p-1]} + b^{[p]})$

Para x_1 :

$$z_1^{[1]} = w^{[1]} x_1^{[0]} + b^{[1]} = [5 \ 6 \ 5]^T$$

$$x_1^{[1]} = \phi^{[1]} \cdot (z_1^{[1]}) = f(z_1^{[1]}) = [0,46212 \ 0,76159 \ 0,46212]^T$$

Da mesma forma:

$$z_1^{[2]} = [4,97061 \ 2,68583]^T$$

$$x_1^{[2]} = [0,45048 \ -0,57642]^T$$

Nota: Neste caso:

$$\phi^{[2]}(z) = \phi(x) = \phi(x^{[3]}) = f(x)$$

$$z_1^{[3]} = [0,87406 \ 1,77503 \ 0,87406]^T$$

$$x_1^{[3]} = [-0,91590 \ -0,80494 \ -0,91590]^T$$

③

De seguida calculamos os deltas para x_1 :

$$\delta \text{ da última camada: } \delta^{[P]} = (x_i^{[P]} - t_i) \circ \phi'^{[P]}(z^{[1]})$$

$$\delta \text{ das restantes camadas: } \delta^{[p]} = w^{[p+1]^T} \cdot \delta_i^{[p+1]} \circ \phi'^{[p]}(z_i^{[p]})$$

$$\frac{d}{dx} \tanh(x) = 1 - \tanh^2(x) \Rightarrow \frac{d}{dx} \tanh(0,5x-2) = \frac{1}{2} (1 - \tanh^2(0,5x-2))$$

$$\text{Portanto } \phi'^{[P]} = f'(x) = \frac{1}{2} [1 - \tanh^2(0,5x-2)]$$

$$\text{Assim, } \delta_1^{[3]} = (x_1^{[3]} - t_1) \circ f'(z_1^{[3]}) = [0,00678 \quad -0,31773 \quad 0,00678]^T$$

$$\delta^{[2]} = w^{[3]^T} \cdot \delta_1^{[3]} \circ f'(z_1^{[2]}) = [-0,37448 \quad -0,110156]^T$$

$$\text{Deste modo, } \delta^{[1]} = [-0,18719 \quad -0,33587 \quad -0,18719]^T$$

Para x_2 executamos exatamente o mesmo algoritmo e obtenos os seguintes resultados:

Forward Propagation:

$$z_2^{[1]} = [1 \ 1 \ 1]^T$$

$$x_2^{[1]} = [-0,90515 \quad -0,90515 \quad -0,90515]^T \quad \delta_2^{[2]} = [-1,15093 \cdot 10^{-5} \quad -1,72899 \cdot 10^{-4}]^T$$

$$z_2^{[2]} = [-4,43089 \quad -1,71544]^T$$

$$\delta_2^{[1]} = [-1,66619 \cdot 10^{-5} \quad -1,97816 \cdot 10^{-5} \quad -1,66619 \cdot 10^{-5}]^T$$

$$x_2^{[2]} = [-0,99936 \quad -0,99343]^T$$

$$z_2^{[3]} = [-0,99294 \quad -2,99212 \quad -0,99299]^T$$

$$x_2^{[3]} = [-0,98652 \quad -0,99816 \quad -0,98652]^T$$

Calculo de deltas:

$$\delta_2^{[3]} = [-2,65961 \cdot 10^{-2} \quad 3,36962 \cdot 10^{-6} \quad 1,80463 \cdot 10^{-5}]^T$$

$$\text{Backpropagation: } w^{[p]} = w^{[p]} - \eta \left(\frac{\partial E^{(1)}}{\partial w^{(p)}} + \frac{\partial E^{(2)}}{\partial w^{(p)}} \right) =$$

$$= w^{[p]} - \eta \cdot \left[\delta_1^{[p]} \cdot (x_1^{[p-1]})^T + \delta_2^{[p]} \cdot (x_2^{[p-1]})^T \right]$$

$$w^{[3]} = \begin{bmatrix} 0,99704 & 0,99775 \\ 3,01431 & 0,98169 \\ 0,99971 & 1,00041 \end{bmatrix}$$

$$w^{[2]} = \begin{bmatrix} 1,01730 & 1,02852 & 1,01730 \\ 1,00468 & 1,00772 & 1,00468 \end{bmatrix}$$

$$w^{[1]} = \begin{bmatrix} 1,01872 & 1,01872 & 1,01872 & 1,01872 \\ 1,03359 & 1,03359 & 2,03359 & 1,03359 \\ 1,01872 & 1,01872 & 1,01872 & 1,01872 \end{bmatrix}$$

$$b^{[p]} = b^{[p]} - \eta \cdot \left[\frac{\partial E^{(1)}}{\partial b^{(p)}} + \frac{\partial E^{(2)}}{\partial b^{(p)}} \right] = b^{[p]} - \eta \cdot [\delta_1^{[p]} + \delta_2^{[p]}]$$

$$b^{[3]} = [1,00191 \quad 1,03177 \quad 0,99930]^T$$

$$b^{[2]} = [1,03745 \quad 1,01017]^T$$

$$b^{[1]} = [1,01872 \quad 1,03359 \quad 1,01872]^T$$

G11_notebook

October 20, 2023

1 Homework 2

Gonçalo Meneses, 103401 e Tomás Arêde, 103239

Antes de começar, vamos remover os avisos que surgem ao longo deste notebook.

```
[ ]: import warnings  
warnings.filterwarnings('ignore')
```

Começamos, então, por carregar o nosso ficheiro *csv* e criar o data frame que irá armazenar os nossos dados.

```
[ ]: import pandas as pd  
from sklearn.model_selection import train_test_split  
  
data = pd.read_csv("winequality-red.csv", delimiter=';')  
  
df = pd.DataFrame(data)  
  
X = df.drop('quality', axis=1)  
y = df['quality']  
  
df.head()
```

```
[ ]:   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4          0.70         0.00          1.9       0.076
1           7.8          0.88         0.00          2.6       0.098
2           7.8          0.76         0.04          2.3       0.092
3          11.2          0.28         0.56          1.9       0.075
4           7.4          0.70         0.00          1.9       0.076

      free sulfur dioxide  total sulfur dioxide  density     pH  sulphates \
0                  11.0            34.0    0.9978  3.51      0.56
1                  25.0            67.0    0.9968  3.20      0.68
2                  15.0            54.0    0.9970  3.26      0.65
3                  17.0            60.0    0.9980  3.16      0.58
4                  11.0            34.0    0.9978  3.51      0.56

alcohol  quality
```

0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

1.0.1 Exercício 1

```
[ ]: from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from IPython.display import set_matplotlib_formats
set_matplotlib_formats('svg', 'pdf')

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,
random_state=0)

residue_list = []

mae_scores = []
mae_rounded_scores = []

rmse_early_scores = []

for state in range(1, 11):

    mlp_regressor = MLPRegressor(hidden_layer_sizes=(10, 10),
activation='relu', early_stopping=True, validation_fraction=0.2,
random_state=state)
    mlp_regressor.fit(X_train, y_train)

    y_pred = mlp_regressor.predict(X_test)

    residues = np.abs(y_pred - y_test)
    residue_list.extend(residues)

    mae = mean_absolute_error(y_true=y_test, y_pred=y_pred)
    mae_scores.append(mae)

    y_pred_rounded = np.round(y_pred)
    y_pred_rounded_bounded = np.clip(y_pred_rounded, 1, 10)

    mae_rounded = mean_absolute_error(y_true=y_test,
y_pred=y_pred_rounded_bounded)
```

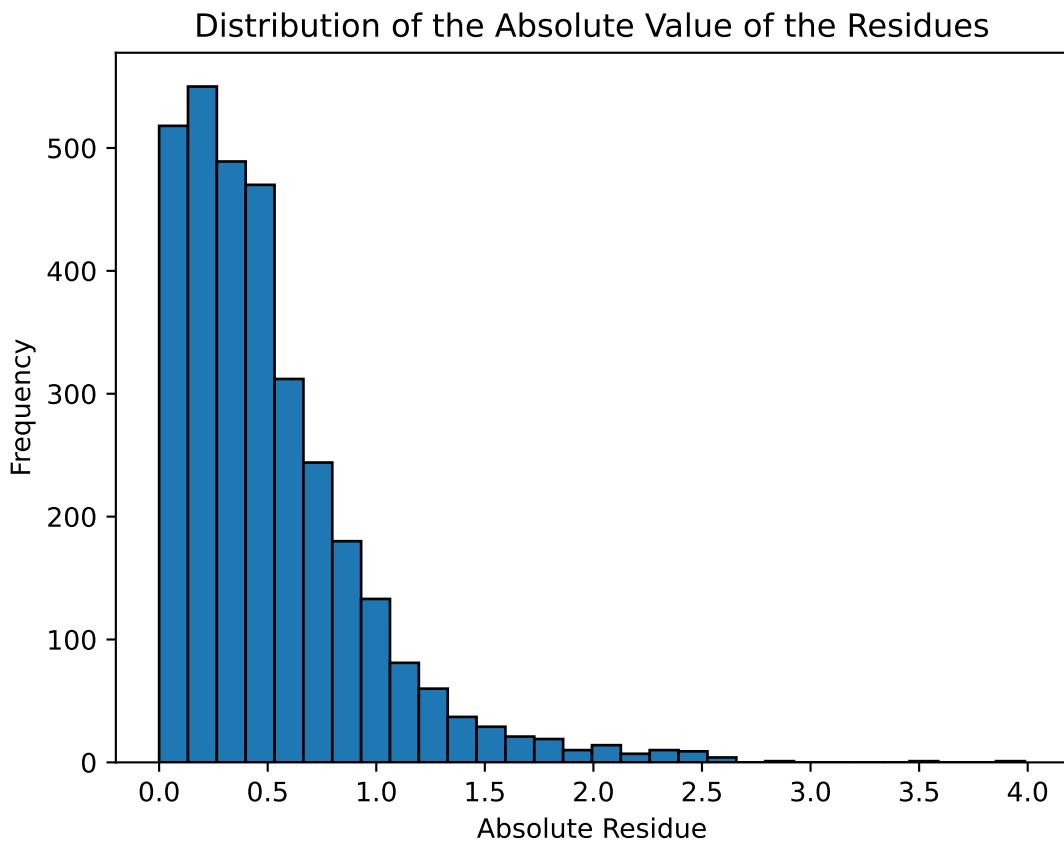
```

mae_rounded_scores.append(mae_rounded)

rmse_early = np.sqrt(mean_squared_error(y_true=y_test, y_pred=y_pred))
rmse_early_scores.append(rmse_early)

plt.hist(residue_list, bins=30, edgecolor='k')
plt.xlabel('Absolute Residue')
plt.ylabel('Frequency')
plt.title('Distribution of the Absolute Value of the Residues')
plt.show()

```



1.0.2 Exercício 2

```

[ ]: avg_mae_scores = np.average(mae_scores)
avg_rounded_mae_scores = np.average(mae_rounded_scores)

print(f'Average of the Mean Absolute Errors: {avg_mae_scores:.5f}')
print(f'Average of the Mean Absolute Erros with rounding and bounding: {avg_rounded_mae_scores:.5f}')

```

```

if avg_rounded_mae_scores < avg_mae_scores:
    print('Rounding and bounding estimates is beneficial to the performance of our MLP classifier.')
else:
    print("Our MLP classifier doesn't benefit from rounding and bounding estimates.")

```

Average of the Mean Absolute Errors: 0.50972
 Average of the Mean Absolute Erros with rounding and bounding: 0.43875
 Rounding and bounding estimates is beneficial to the performance of our MLP classifier.

1.0.3 Exercício 3

```

[ ]: num_iterations = [20, 50, 100, 200]

rmse_scores = []

for random_state in range(1, 11):
    for iter in num_iterations:

        mlp_regressor = MLPRegressor(hidden_layer_sizes=(10, 10), activation='relu', max_iter=iter, random_state=random_state)

        mlp_regressor.fit(X_train, y_train)

        y_pred = mlp_regressor.predict(X_test)

        rmse = np.sqrt(mean_squared_error(y_true=y_test, y_pred=y_pred))
        rmse_scores.append((iter, rmse))

avg_rmse_early_score = np.average(rmse_early_scores)

print(f'The average RMSE score for Early Stopping is: {avg_rmse_early_score:.5f}')

average_rmse_scores = {}

for iter, rmse in rmse_scores:
    average_rmse_scores.setdefault(iter, []).append(rmse)

iters, avg_rmse_values = zip(*[(iter, np.mean(rmse_list)) for iter, rmse_list in average_rmse_scores.items()])
[print(f'The average RMSE for {iter} iterations is: {avg_rmse:.5f}') for iter, avg_rmse in zip(iters, avg_rmse_values)]

plt.figure(figsize=(8, 5))

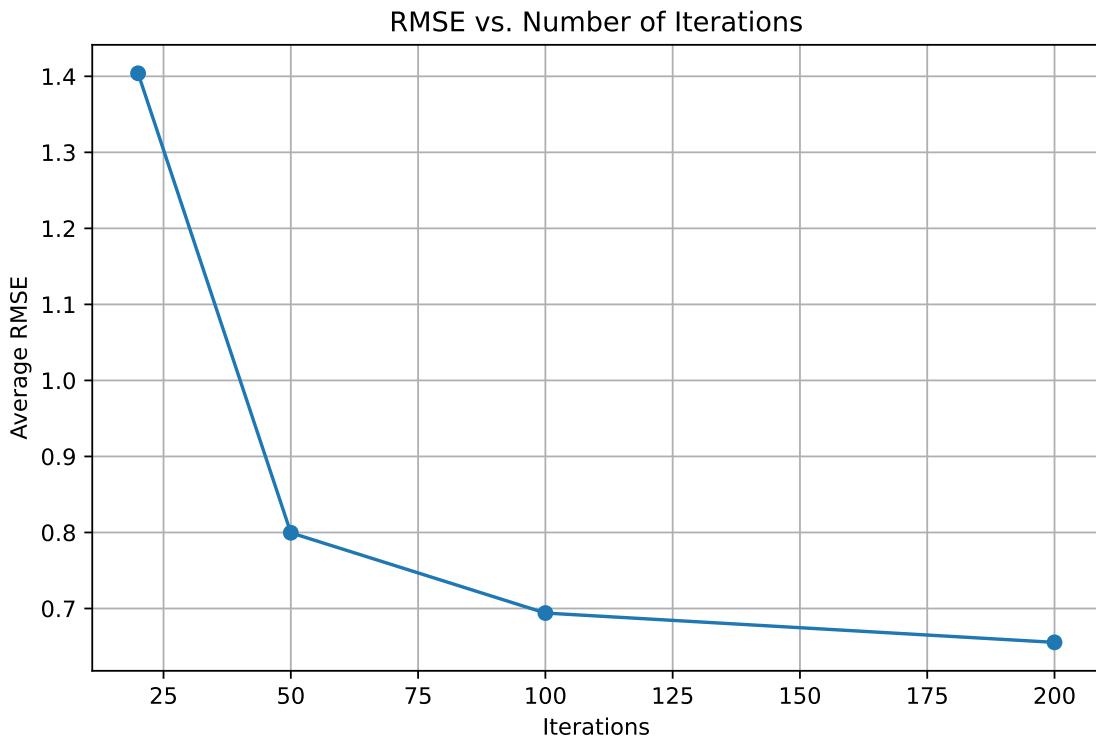
```

```

plt.plot(iters, avg_rmse_values, marker='o')
plt.title('RMSE vs. Number of Iterations')
plt.xlabel('Iterations')
plt.ylabel('Average RMSE')
plt.grid(True)
plt.show()

```

The average RMSE score for Early Stopping is: 0.67065
The average RMSE for 20 iterations is: 1.40398
The average RMSE for 50 iterations is: 0.79961
The average RMSE for 100 iterations is: 0.69404
The average RMSE for 200 iterations is: 0.65545



1.0.4 Exercício 4

Apesar de, como se pode ver acima, o valor do *RMSE* ter diminuído com o número de iterações, é importante destacar alguns pontos.

Primeiro, o valor do erro associado a 20 iterações é de longe o mais alto, quando comparado aos das restantes iterações. Isso significa que, neste caso, o nosso modelo requer um grande número de iterações para treinar o modelo, ou seja, se utilizarmos *Early Stopping*, este poderá causar um término precoce do nosso treino, o que levará a um *underfitting* e, consequentemente, a uma performance subóptima.

Segundo, à medida que o número de iterações aumenta (neste caso, 20, 50, 100, 200), a performance

do nosso modelo também, e, em troca, o valor do *RMSE* diminui, porém, em incrementos cada vez mais pequenos. Assim, para um número suficientemente grande de iterações poderá haver um *overfitting* para os dados de treino, pelo que a utilização de *Early Stopping*, face ao número máximo de iterações, impediria que tal acontecesse. Isto é possível, uma vez que este método de regularização para o treino quando a performance do modelo para um conjunto de validação diminui e mantém a sua capacidade de generalização.

Finalmente, podemos concluir que um bom balanço entre *Early Stopping* e um número adequado de iterações é essencial para uma performance óptima do modelo, impedindo tanto um *underfit* como um *overfit*.